

AutoJudge: Automatic Programming Problem Difficulty Prediction

1. Introduction

Online competitive programming platforms classify problems into difficulty levels such as *Easy*, *Medium*, and *Hard*, and often assign a numerical difficulty score. These labels are usually determined manually using expert judgment and user feedback, which is subjective and time-consuming.

This project, **AutoJudge**, aims to automate this process using Natural Language Processing (NLP) and Machine Learning. Given only the **textual description** of a programming problem, the system predicts:

1. **Problem Class** – Easy / Medium / Hard (classification)
2. **Problem Score** – a continuous numerical difficulty value (regression)

The project also includes a simple **web interface** where users can paste a new problem statement and obtain predictions in real time.

2. Dataset Description

The dataset used in this project was **provided as part of the problem statement itself**. No external data sources were collected or labeled by the authors.

Each data sample contains the following fields:

- `title`
- `description`
- `input_description`
- `output_description`
- `problem_class` (Easy / Medium / Hard)
- `problem_score` (numerical difficulty score)

All labels (`problem_class` and `problem_score`) were already included in the dataset, as explicitly specified in the official project problem statement provided by the club. The task was strictly limited to **using the provided dataset only**, without any manual relabeling or data augmentation.

For modeling purposes, the textual fields (`description`, `input_description`, and `output_description`) were concatenated into a single field called **combined_text**, following the official project instructions.

3. Data Preprocessing

The following preprocessing steps were applied:

1. Lowercasing all text
2. Removal of English stop words
3. Tokenization handled internally by TF-IDF
4. Numerical constraints and symbols retained as they convey problem complexity

No data leakage was introduced; all transformations were fitted only on training data.

4. Feature Engineering

4.1 TF-IDF Features

A **TF-IDF Vectorizer** was used to capture semantic and syntactic patterns from problem descriptions.

- Max features: 12,000
- N-grams: (1, 2)
- Stop words: English

4.2 Handcrafted Complexity Features

To complement TF-IDF, domain-specific features were engineered:

- Text length
- Number of mathematical symbols
- Frequency of algorithmic keywords (DP, graph, BFS, DFS, etc.)
- Average word length
- Sentence count
- Presence of explicit algorithm hints
- Maximum numerical constraint
- Control-flow keyword density (`if`, `for`, `while`)

These features capture **structural and logical complexity** beyond raw text semantics.

4.3 Feature Scaling and Selection

- Min-Max scaling was applied to handcrafted features (chi-square safe)
- TF-IDF and handcrafted features were concatenated
- **Chi-square feature selection** reduced features to the top 5000

Final feature shape: **(4112, 5000)**

5. Classification Model

5.1 Motivation for Hierarchical Classification

A flat 3-class classifier struggled with the *Easy* class due to overlap with Medium problems. To address this, a **two-stage hierarchical approach** was adopted.

5.2 Stage 1: Easy vs Non-Easy

- Model: Linear SVM (LinearSVC)
- Class weights: Easy boosted ($2.5\times$)
- Objective: Maximize recall for Easy problems

Stage-1 Accuracy: 80.80%

5.3 Stage 2: Medium vs Hard

- Model: Linear SVM (LinearSVC)
- Balanced class weights
- Trained only on non-easy samples

5.4 Final Prediction Logic

1. Predict Easy vs Non-Easy
 2. If Non-Easy → predict Medium or Hard
-

6. Classification Results

Overall Accuracy

Final Hierarchical Accuracy: 79.95%

Classification Report

- Easy: Precision 0.67, Recall 0.41
- Medium: Precision 0.79, Recall 0.88
- Hard: Precision 0.84, Recall 0.89

The hierarchical approach significantly improved robustness compared to a flat classifier, especially for Medium and Hard classes.

7. Regression Model (Difficulty Score Prediction)

7.1 Motivation

Difficulty scores vary significantly across classes. A single regressor underperformed due to heterogeneity. Hence, **class-conditional regression** was adopted.

7.2 Model

- Separate **Linear SVR** models for Easy, Medium, and Hard
- Shared feature space from classification

7.3 Evaluation Metrics

- **MAE:** 0.795
- **RMSE:** 1.017
- **R² Score:** 0.787

This demonstrates strong predictive performance without overfitting or data leakage.

8. Web Interface

A lightweight web interface was developed using **Streamlit** and deployed publicly for live demonstration and verification.

- **Live Application URL:** <https://autojudge-hxoigre646axmf2kpmdtbi.streamlit.app/>

User Input Format

The interface strictly follows the project specification:

- **Problem Description** – Main problem statement
- **Input Description** – Description of input format and constraints
- **Output Description** – Description of expected output

The user fills all three text boxes and clicks **Predict**.

Backend Processing

1. All three inputs are concatenated internally into a single text string
2. TF-IDF and handcrafted features are extracted
3. Hierarchical SVM predicts the difficulty class
4. The corresponding class-conditional SVR predicts the difficulty score

Output Displayed

- Predicted difficulty class (Easy / Medium / Hard)

- Predicted numerical difficulty score

This deployed interface loads the same serialized models (.pk1 files) submitted in the GitHub repository, ensuring full consistency between reported results, live predictions, and the submitted code.

9. Conclusion

This project demonstrates that programming problem difficulty can be effectively predicted using only textual descriptions. Key contributions include:

- Hybrid TF-IDF + handcrafted feature engineering
- Cost-sensitive hierarchical classification
- Class-conditional regression for robust score prediction

Future improvements may include transformer-based embeddings and cross-platform generalization.

10. Repository Structure

The submitted GitHub repository contains **all necessary executable code and trained models** required to verify the project results.

Current repository contents: - `app.py` - Streamlit web application (loads trained models and performs inference) - `tfidf.pk1` - Trained TF-IDF vectorizer - `extra_feature_scaler.pk1` - MinMax scaler for handcrafted features - `selector.pk1` - Chi-square feature selector (top 5000 features) - `stage1_easy_vs_noneasy.pk1` - Stage-1 hierarchical classifier (Easy vs Non-Easy) - `stage2_medium_vs_hard.pk1` - Stage-2 classifier (Medium vs Hard) - `svr_models.pk1` - Class-conditional regression models - `requirements.txt` - Python dependencies - `README.md` - Project overview and usage instructions - `AutoJudge_Training_Reference.ipynb` - Reference Colab notebook

Note on Training Code

The primary training experiments for this project were conducted offline, and the final trained models are stored as serialized .pk1 files in the repository. These trained artifacts are the ones used by the web application and reported in this document.

For additional clarity and transparency, a **reference Colab notebook** (`AutoJudge_Training_Reference.ipynb`) has been included in the repository. This notebook reproduces the same training pipeline and methodology and is provided **only for reference and reproducibility**. The reported results in this document correspond to the trained models and not to any subsequent retraining.

This setup ensures:

- Consistency between reported results and deployed models
- Reproducibility for reviewers
- Faster verification without requiring model retraining